

```
In [2]: import pandas as pd
import re
import nltk
import matplotlib.pyplot as plt
import seaborn as sns
from nltk.corpus import stopwords
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score, confusion_mat
from textblob import TextBlob # Simple sentiment analysis tool

# Download NLTK stopwords
nltk.download('stopwords')
stop_words = set(stopwords.words('english'))
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\sujal\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
In [3]: # ----- Load and Combine Datasets ----- #
print("\nLoading datasets...")
df1 = pd.read_csv("tweets/data_analysis.csv", low_memory=False)
df2 = pd.read_csv("tweets/data_science.csv", low_memory=False)
df3 = pd.read_csv("tweets/data_visualization.csv", low_memory=False)

df = pd.concat([df1, df2, df3], ignore_index=True)
print("✅ Datasets loaded and combined. Total rows:", df.shape[0])
```

Loading datasets...

✅ Datasets loaded and combined. Total rows: 314092

```
In [4]: # ----- Inspect Columns ----- #
print("\nAvailable columns:", df.columns.tolist())

# Rename 'tweet' to 'text' if available
if 'tweet' in df.columns:
    df.rename(columns={'tweet': 'text'}, inplace=True)
    print("✅ Renamed 'tweet' column to 'text'.")
```

Available columns: ['id', 'conversation\_id', 'created\_at', 'date', 'time', 'timezone', 'user\_id', 'username', 'name', 'place', 'tweet', 'language', 'mentions', 'urls', 'photos', 'replies\_count', 'retweets\_count', 'likes\_count', 'hashtags', 'cashtags', 'link', 'retweet', 'quote\_url', 'video', 'thumbnail', 'near', 'geo', 'source', 'user\_rt\_id', 'user\_rt', 'retweet\_id', 'reply\_to', 'retweet\_date', 'translate', 'trans\_src', 'trans\_dest']

✅ Renamed 'tweet' column to 'text'.

```
In [5]: # ----- Assign Sentiment Labels ----- #
print("\nAssigning sentiment using TextBlob...")
df['sentiment'] = df['text'].apply(lambda x: 'positive' if TextBlob(str(x)).sentiment.polarity > 0 else 'negative')
print("✅ Sentiment assigned based on polarity.")

# Show sample rows
print("\nSample labeled data:")
print(df[['text', 'sentiment']].head(5))
```

Assigning sentiment using TextBlob...

✔ Sentiment assigned based on polarity.

Sample labeled data:

	text	sentiment
0	Join @SPJ_IC at 6 p.m. EDT Tuesday for the nex...	negative
1	READ: Resilience, hard work, and perseverance ...	negative
2	How will the shift to streaming impact movie-m...	positive
3	(2/2) to gain the aptitude to find better solu...	positive
4	Alexander Lukashenko is still in power in Bela...	negative

```
In [7]: # ----- Clean and Filter Data ----- #
print("\nCleaning text...")
df = df[['text', 'sentiment']].dropna()

def clean_text(text):
    text = text.lower()
    text = re.sub(r"http\S+|www\S+|https\S+", '', text)
    text = re.sub(r'\@w+|\#', '', text)
    text = re.sub(r'^A-Za-z\s', '', text)
    tokens = text.split()
    tokens = [word for word in tokens if word not in stop_words]
    return " ".join(tokens)

df['clean_text'] = df['text'].apply(clean_text)
print("✔ Text cleaned.")

# Show cleaned sample
print("\nSample cleaned text:")
print(df[['clean_text']].head(5))
```

Cleaning text...

✔ Text cleaned.

Sample cleaned text:

	clean_text
0	join pm edt tuesday next ictalk created code b...
1	read resilience hard work perseverance paid ge...
2	shift streaming impact moviemaking interesting...
3	gain aptitude find better solutions realworld ...
4	alexander lukashenko still power belarus despi...

```
In [8]: # ----- Encode Labels and Split ----- #
print("\nEncoding sentiment labels...")
df['label'] = df['sentiment'].map({'positive': 1, 'negative': 0})
print("✔ Labels encoded: 1 for Positive, 0 for Negative")

print("\nSentiment label distribution:")
print(df['label'].value_counts())

print("\nSplitting data into train and test sets (80/20)...")
X_train, X_test, y_train, y_test = train_test_split(
    df['clean_text'], df['label'], test_size=0.2, random_state=42
)
print(f"✔ Training samples: {len(X_train)}, Testing samples: {len(X_test)}")
```

Encoding sentiment labels...

✅ Labels encoded: 1 for Positive, 0 for Negative

Sentiment label distribution:

label

0 164938

1 149154

Name: count, dtype: int64

Splitting data into train and test sets (80/20)...

✅ Training samples: 251273, Testing samples: 62819

```
In [9]: # ----- Vectorization ----- #
print("\nVectorizing text...")
vectorizer = CountVectorizer()
X_train_vec = vectorizer.fit_transform(X_train)
X_test_vec = vectorizer.transform(X_test)
print("✅ Text vectorized. Feature count:", X_train_vec.shape[1])
```

Vectorizing text...

✅ Text vectorized. Feature count: 120074

```
In [10]: # ----- Train Model ----- #
print("\nTraining Logistic Regression model...")
model = LogisticRegression()
model.fit(X_train_vec, y_train)
print("✅ Model training complete.")
```

Training Logistic Regression model...

✅ Model training complete.

c:\Users\sujal\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

```
In [12]: # ----- Evaluate Model ----- #
print("\nPredicting on test data...")
y_pred = model.predict(X_test_vec)

print("\n--- Model Evaluation ---")
print("Accuracy Score:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Predicting on test data...

--- Model Evaluation ---

Accuracy Score: 0.9444753975707987

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.97	0.95	33050
1	0.96	0.92	0.94	29769
accuracy			0.94	62819
macro avg	0.95	0.94	0.94	62819
weighted avg	0.95	0.94	0.94	62819

```
In [13]: # ----- Confusion Matrix ----- #
print("\nGenerating confusion matrix...")
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", cm)


sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Negative', 'Positive'],
            yticklabels=['Negative', 'Positive'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.tight_layout()
plt.savefig("confusion_matrix.png")
plt.show()

print("✅ Done. Confusion matrix saved as 'confusion_matrix.png'")
```

Generating confusion matrix...

Confusion Matrix:

```
[[31941 1109]
 [ 2379 27390]]
```



✅ Done. Confusion matrix saved as 'confusion\_matrix.png'